



Relax-and-Recover



Relax-and-Recover (ReaR) Automated Testing

Gratien D'haese
IT3 Consultants

<https://gdha.github.io/rear-automated-testing/>



Agenda

- Who am I?
- What is Relax-and-Recover?
- Relax-and-Recover Automated Testing
- DEMO



Who am I?



- Gratiën D'haese
- IT3 Consultants (company)
 - > 25 years Unix experience
 - Unix/Linux Engineer (incl. DevOps)
 - ReaR Support Contracts
- Relax-and-Recover (ReaR)
 - Major Open Source project
 - <https://github.com/gdha>



Circle of Wisdom





Linux Disaster Recovery

Question: “What shall I do if a disaster strikes?”


Dependent on:

- Hardware failure (e.g. boot disk lost)
- Lost everything (fire, water, earthquake, theft)
- The answer: “Act immediately (with a disaster recovery plan)”




Why are backups not enough?

Disaster Recovery via re-install of Linux OS and restore data



Disaster Recovery with ReaR and restore data



Time





Relax-and-Recover (ReaR) as DR solution

- ReaR is a tool that implements a **DR work-flow**
- Basically meaning:
 - Modular framework written in **Bash**
 - Easy to extend to own needs
 - Easy to deploy (set up and forget)
 - Integration for various Linux technologies
 - Integration with various back-up solutions
 - Attempts to make system recovery *as easy as possible*



Disaster Recovery – How It Works

- Store the disk layout
 - Partitioning, LVM and RAID configuration
 - File systems, file system labels ...
 - Boot loader (GRUB, GRUB2, LILO, UEFI)
- Store the files (tgz, rsync, through backup software ...)
- Create bootable rescue media with system configuration (and backup data)
- **Can be done online** (no business interruption)




Disaster Recovery – Rescue Media

- Create “rescue linux” from running system
- Optimally compatible “tool box”
- Clone the system environment
 - Linux kernel and modules
 - Device driver configuration
 - Network configuration
 - Basic system software and tools
- Operate entirely in RAM (initrd)





DR Flow – BACKUP and OUTPUT





Usage of rear

- Shell scripts are stored under `/usr/share/rear`
- Scripts are executed via work-flows:
 - **mkrescue** (only make rescue image)
 - **mkbackup** (including make rescue image)
 - **mkbackuponly** (excluding make rescue image)
 - **recover** (the actual recovery part)
- Easy to incorporate new scripts, e.g. for information gathering of Hard- and Software, or other goodies



Getting started with ReaR

- Download it from
 - Our web-site
 - <http://relax-and-recover.org/download/>
 - The rear-snapshot rpm's build from GitHub
 - <http://download.opensuse.org/repositories/Archiving:/Backup:/Rear:/Snapshot/>
 - The official source
 - <https://github.com/rear/rear>
 - The official repo's (Fedora, RHEL and SLES)
 - yum install rear
 - zypper install rear



Testing ReaR

- ReaR is due to a wide range of options difficult to test
- Different Linux flavors are using similar tools with some minor differences :-(
- Too much to be able to test everything before a new release, and sometimes it is even not possible just because we do not have the hardware
- Is and will stay a challenge for the future...




ReaR Builds

- Finding a right balance between CI Testing and Automated ReaR Testing
- Via **OpenSuse Build Services** we build daily a fresh ReaR snapshot package (for free)
- We have some excellent ReaR developers who do lots of coding (for free)
- Support via GitHub issues
 - Free support
 - Commercial support



ReaR Automated Testing Workflow





Automated ReaR Testing

- Wrote it for customers with a ReaR support or subscription contract
- Currently we support the following GNU/Linux distributions:
 - CentOS 7
 - Ubuntu 14.04, Ubuntu 16.04
 - SLES 11, SLES 12
- The VMs are provisioned with ansible playbooks




Automated ReaR Testing (cont'd)

- We start the process and it automatically does
 - DR image creation
 - Backup of system via:
 - BACKUP=NETFS (with GNU tar)
 - BACKUP=DUPLICITY
 - BACKUP=BAREOS
 - PXE or ISO boot the recover system with DR image made
 - Restore backup
 - Reboot the recover system



Test Configuration





Set up vagrant environment

- Host system must be GNU/Linux, or Mac OS based
- A hypervisor like VirtualBox (or KVM on Linux)
- Install “vagrant” from your distro, or from <https://www.vagrantup.com/downloads.html>
- KVM with libvirt needs the vagrant-libvirt plugin
`vagrant plugin install vagrant-libvirt`
- Install “git” software to clone the Vagrantfile and scripts



Install the ReaR automated Testing software

- Is Open Source and licensed under GPLv3
- New code is written only for customers with a valid ReaR Support contract (*PR are welcome*)
- `git clone git@github.com:gdha/rear-automated-testing.git`
- Go into directory “rear-automated-testing”
- Type “./rear-automated-test.sh -h” to see info
- Uses “*vagrant*” to drive the creation of the VMs
- Account vagrant/vagrant (and root/vagrant)




Login via vagrant or ssh

- Once the VMs are up and running
- Login via vagrant:
 - `sudo vagrant ssh client`
 - `sudo vagrant ssh server`
- Another way to login is via ssh:
 - `ssh root@192.168.33.10` (client root pw is vagrant)
 - `ssh root@192.168.33.15` (server root pw is vagrant)



Login via vncviewer

- If you install “tigervnc” you can use **vncviewer**
- Use address 127.0.0.1 (localhost)
- Port 5991 for “client”
- Port 5992 for “server”
- Port 5993 for “recover”





Try it yourself?

- <https://github.com/gdha/rear-automated-testing>

```
$ sudo ./rear-automated-test.sh -p libvirt
```

```
+-----+
| Relax-and-Recover Automated Testing script |
|               version 1.3                  |
+-----+
```

Author: Gratien D'haese

Copyright: GPL v3

Current distro directory is centos7

Copy the Vagrantfile.libvirt to Vagrantfile

Bringing up the vagrant VMs client and server

Bringing machine 'client' up with 'libvirt' provider...

Bringing machine 'server' up with 'libvirt' provider...

==> server: Starting domain.

==> client: Starting domain.



ReaR Automated Testing

- ReaR Automated Testing speeds up
 - Validation processes
 - Bug hunting
 - Compliance checks
 - Trusworthy of “unstable” ReaR repository
 - Acceptance of ReaR within corporations and Open Source vendors
 - Stable releases can be tested anytime
- We accept ***pull requests*** and sponsoring

ReaR Automated Testing

Demonstration



Need more of this?

We can foresee in a customized workshop on consultancy basis, or set-up in-house full automated ReaR testing for customers with a valid support contract

See <http://www.it3.be/rear-support>



Gratien D'haese
IT3 Consultants



Backup Slides in case live demonstration
is not possible



The help and usage page

```
$ sudo ./rear-automated-test.sh -h
```

Usage: rear-automated-test.sh [-d distro] [-b <boot method>] [-s <stable rear version>] [-p provider] [-c rear-config-file.conf] [-t test] -vh

-d: The distribution to use for this automated test (default: centos7)

-b: The boot method to use by our automated test (default: PXE)

-s: The <stable rear version> is the specific version we want to test, e.g. 2.3 (default: <empty>)

-p: The vagrant <provider> to use (default: virtualbox)

-c: The ReaR config file we want to use with this test (default: PXE-booting-with-URL-style.conf)

-l: The ReaR test logs top directory (default: /export/rear-tests/logs)

-t: The ReaR validation test directory (see tests directory; no default)

-h: This help message.

-v: Revision number of this script.



Starting the script

```
$ sudo ./rear-automated-test.sh
```

```
+-----+
| Relax-and-Recover Automated Testing script |
|           version 1.0                      |
+-----+
```

Author: Gratien D'haese

Copyright: GPL v3

Current distro directory is centos7

Copy the Vagrantfile.virtualbox to Vagrantfile

Bringing up the vagrant VMs client and server

Bringing machine 'client' up with 'virtualbox' provider...

Bringing machine 'server' up with 'virtualbox' provider...

==> client: Checking if box 'centos/7' is up to date...

==> client: A newer version of the box 'centos/7' is available! You currently



```
==> server: Setting hostname...
==> server: Configuring and enabling network interfaces...
==> server: Rsyncing folder: /projects/rear/rear-automated-testing/centos7/ => /vagrant
==> server: Machine already provisioned. Run `vagrant provision` or use the `--provision`
==> server: flag to force provisioning. Provisioners marked to run always will still run.
```

Sleep for 5 seconds [Control-C is now possible]

Current machine states:

client	running (virtualbox)
server	running (virtualbox)
recover	poweroff (virtualbox)

This environment represents multiple VMs. The VMs are all listed above with their current state. For more information about a specific VM, run `vagrant status NAME`.



Check if 'eth1' is active on client [known issue
<https://github.com/mitchellh/vagrant/issues/8166>]
Connection to 127.0.0.1 closed.
Check if 'eth1' is active on server
Connection to 127.0.0.1 closed.
Doing ping tests to VMs client and server
client is up and running - ping test **OK**
server is up and running - ping test **OK**



Update rear on the VM client

Loaded plugins: fastestmirror

Determining fastest mirrors

* base: ftp.belnet.be

Resolving Dependencies

--> Running transaction check

---> Package rear.x86_64 0:2.00-44.git201702131649.el7 will be updated

---> Package rear.x86_64 0:2.00-74.git201703170925.el7 will be an update

--> Finished Dependency Resolution

Dependencies Resolved

=====				
Package	Arch	Version	Repository	Size
=====				

Updating:

rear x86_64 2.00-74.git201703170925.el7 Archiving_Backup_Rear_Snapshot 459 k



ReaR version that will be tested is:

Relax-and-Recover 2.00-git201703170925 / 2017-03-17

Content of /etc/rear/local.conf is:

OUTPUT=PXE

OUTPUT_PREFIX_PXE=\$HOSTNAME

BACKUP=NETFS

BACKUP_URL=nfs://192.168.33.15/export/nfs/tftpboot

PXE_TFTP_URL=nfs://10.0.2.2/root/.config/VirtualBox/TFTP

PXE_CONFIG_URL=nfs://10.0.2.2/root/.config/VirtualBox/TFTP/pxelinux.cfg

PXE_RECOVER_MODE="unattended"

PXE_CREATE_LINKS="IP"

USE_STATIC_NETWORKING=y

KERNEL_CMDLINE="\$KERNEL_CMDLINE net.ifnames=0"



Run rear -v mkbakup

Relax-and-Recover 2.00-git201703170925 / 2017-03-17

Using log file: /var/log/rear/rear-client.log

Using backup archive '/tmp/rear.mbxNVLH8555SO9t/outputfs/client/backup.tar.gz'

Creating disk layout

Creating root filesystem layout

Copying logfile /var/log/rear/rear-client.log into initramfs as '/tmp/rear-client-partial-2017-03-19T09:43:41+0100.log'

Copied kernel+initrd 108M to nfs://10.0.2.2/root/.config/VirtualBox/TFTP

Created pxelinux config 'rear-client' and symlinks for IP addresses in
nfs://10.0.2.2/root/.config/VirtualBox/TFTP/pxelinux.cfg

Copying resulting files to nfs location

Saving /var/log/rear/rear-client.log as rear-client.log to nfs location

Creating tar archive '/tmp/rear.mbxNVLH8555SO9t/outputfs/client/backup.tar.gz'

Preparing archive operation

Archived 410 MiB in 64 seconds [avg 6567 KiB/sec]

The rear mkbakup was successful



Make client area readable for others on PXE boot server 10.0.2.2

Halting the client VM before doing the recovery

Recover VM will use the client IP address after it has been fully restored

=> client: Attempting graceful shutdown of VM...

Copied private key of client VB to recover VB config area

Starting the recover VM

Bringing machine 'recover' up with 'virtualbox' provider...

=> recover: Checking if box 'clink15/pxe' is up to date...

File Machine View Input Devices Help

Relax-and-Recover 2.00-git201703221726 / 2017-03-22

Relax-and-Recover comes with ABSOLUTELY NO WARRANTY; for details see
the GNU General Public License at: <http://www.gnu.org/licenses/gpl.html>

Host client.box using Backup NETFS and Output PXE
Build date: Thu, 23 Mar 2017 09:33:12 +0100

rear-unattended - Recover client with unattended kernel option

local - Boot from next boot device

boothd0 - boot first local disk

boothd1 - boot second local disk

hdt - Hardware Detection Tool

reboot - Reboot the system

poweroff - Poweroff the system

Loading client/client.kernel.....

Loading client/client.initrd.cgz.....

.....

File Machine View Input Devices Help

```
[ 0.237501] usbcore: registered new interface driver hub
[ 0.240418] usbcore: registered new device driver usb
[ 0.241148] PCI: Using ACPI for IRQ routing
[ 0.241788] NetLabel: Initializing
[ 0.242331] NetLabel: domain hash size = 128
[ 0.242836] NetLabel: protocols = UNLABELED CIPSOv4
[ 0.243120] NetLabel: unlabeled traffic allowed by default
[ 0.243656] Switched to clocksource refined-jiffies
[ 0.247567] pnp: PnP ACPI init
[ 0.248076] ACPI: bus type PNP registered
[ 0.249059] pnp: PnP ACPI: found 2 devices
[ 0.249578] ACPI: bus type PNP unregistered
[ 0.255960] Switched to clocksource acpi_pm
[ 0.256607] NET: Registered protocol family 2
[ 0.257134] TCP established hash table entries: 8192 (order: 4, 65536 bytes)
[ 0.257747] TCP bind hash table entries: 8192 (order: 5, 131072 bytes)
[ 0.258218] TCP: Hash tables configured (established 8192 bind 8192)
[ 0.258573] TCP: reno registered
[ 0.259023] UDP hash table entries: 512 (order: 2, 16384 bytes)
[ 0.259585] UDP-Lite hash table entries: 512 (order: 2, 16384 bytes)
[ 0.260852] NET: Registered protocol family 1
[ 0.261040] pci 0000:00:00.0: Limiting direct PCI/PCI transfers
[ 0.261638] pci 0000:00:01.0: Activating ISA DMA hang workarounds
[ 0.263031] Unpacking initramfs...
```

File Machine View Input Devices Help

Switching off migration mode due to 'unattended' kernel option

Start system layout restoration.

Creating partitions for disk /dev/sda (msdos)

Creating LVM PV /dev/sda3

0 logical volume(s) in volume group "VolGroup00" now active

Restoring LVM VG VolGroup00

Sleeping 3 seconds to let udev or systemd-udev create their devices...

Creating filesystem of type xfs with mount point / on /dev/mapper/VolGroup00-LogVol100.

/dev/mapper/VolGroup00-LogVol100: 4 bytes were erased at offset 0x00000000 (xfs):
58 46 53 42

Mounting filesystem /

Creating filesystem of type xfs with mount point /boot on /dev/sda2.

/dev/sda2: 4 bytes were erased at offset 0x00000000 (xfs): 58 46 53 42

Mounting filesystem /boot

Creating swap on /dev/mapper/VolGroup00-LogVol101

Disk layout created.

Restoring from '/tmp/rear.lZJ8m1Mh05dKpr0/outputfs/client/backup.tar.gz'...

Restored 1001 MiB [avg. 68337 KiB/sec] OK

Restored 1039 MiB in 16 seconds [avg. 66556 KiB/sec]

Restoring finished.

Restore the Mountpoints (with permissions) from /var/lib/rear/recovery/mountpoint_permissions

Running mkinitrd...

File Machine View Input Devices Help

Mounting filesystem /boot

Creating swap on /dev/mapper/VolGroup00-LogVol01

Disk layout created.

Restoring from '/tmp/rear.lZJ8m1Mh05dKpr0/outputfs/client/backup.tar.gz' ...

Restored 1001 MiB [avg. 68337 KiB/sec] OK

Restored 1039 MiB in 16 seconds [avg. 66556 KiB/sec]

Restoring finished.

Restore the Mountpoints (with permissions) from /var/lib/rear/recovery/mountpoint_permissions

Running mkinitrd...

Updated initrd with new drivers for kernel 3.10.0-514.2.2.el7.x86_64.

Skip installing GRUB Legacy boot loader because GRUB 2 is installed (grub-probe or grub2-probe exist).

Installing GRUB2 boot loader

Updated pxelinux config 'rear-client' to boot from first hard disk at next reboot

Finished recovering your system. You can explore it under '/mnt/local'.

Success !

Rebooting in 30 seconds (Ctrl-C to interrupt)

File Machine View Input Devices Help

```
ipXE (PCI C8:00.0) starting execution...ok  
ipXE initialising devices...ok
```

```
ipXE 1.0.0+ -- Open Source Network Boot Firmware -- http://ipxe.org  
Features: DNS TFTP HTTP PXE PXEXT Menu
```

```
net0: 08:00:27:40:34:54 using 82540em on PCI00:03.0 (open)  
[Link:up, TX:0 TXE:0 RX:0 RXE:0]  
DHCP (net0 08:00:27:40:34:54)...._
```


File Machine View Input Devices Help

```
CentOS Linux 7 (Core)  
Kernel 3.10.0-514.2.2.el7.x86_64 on an x86_64  
  
client login: _
```

File Machine View Input Devices Help

CentOS Linux 7 (Core)

Kernel 3.10.0-514.2.2.el7.x86_64 on an x86_64

client login: root

Password:

Last login: Sun Mar 19 09:42:08 on pts/0

[root@client ~]# w

09:48:14 up 0 min, 1 user, load average: 0.83, 0.19, 0.06

USER	TTY	FROM	LOGIN@	IDLE	JCPU	PCPU	WHAT
root	tty1		09:48	6.00s	0.01s	0.00s	w

[root@client ~]# _



Contact

- ✓ Gratien D'haese
e-mail: gratien.dhaese@it3.be
- ✓ web: <http://www.it3.be>
- ✓ Relax-and-Recover main project site:
<http://relax-and-recover.org/>
- ✓ Relax-and-Recover Sources and Issues:
<https://github.com/rear/rear>
- ✓ Commercial Support:
<http://www.it3.be/rear-support/>